
Capacity or Classifier? What Expansion Triggers Detect in Task-Free Continual Adapter Learning, and When Expansion Pays

Anonymous authors

Abstract

Dynamic-expansion continual learners add an adapter when a trigger—a loss spike, a representation shift, a prospective-forgetting estimate, or the utility of a fresh module—signals that existing capacity is inadequate. We ask what these triggers detect and when expanding is worth it. At the arrival of labels that the current module’s classifier has never trained on, the excess loss depends on the input only through the old labels’ logits, so loss-based triggers and one-step fresh-module utility are *classifier-novelty* detectors. In a pre-registered study (ten text streams, three held out; DistilBERT, a held-out BERT backbone and ViT-B/16), every loss-spike and fresh-utility expansion in class-incremental streams fell at package-level label novelty. We then measure the value of expansion directly against two references, a single adapter that keeps learning and first-session adaptation. When the learner keeps exemplars from which its classifier statistics are refreshed, oracle expansion at the true boundaries costs 7–11 points in class-incremental streams relative to a single adapter, is neutral under domain shift and mixed under concept drift, although it beats first-session adaptation, the reference used in self-expansion ablations. When the learner is exemplar-free, so that a training adapter’s prototypes go stale, the sign flips: [EXT2]. Refreshing stale prototypes without exemplars by semantic drift compensation [EXT4], and the same interaction appears inside the official implementation of a published self-expanding method (SEMA) [SEMA]. Expansion at classifier novelty is therefore a remedy for stale classifier statistics rather than for missing capacity, and its value depends on whether those statistics can be refreshed. We recommend that expansion triggers be evaluated by the measured value of expansion against both reference points, not by boundary precision.

1 Introduction

Parameter-efficient continual learning commonly grows a pool of adapters around a frozen pre-trained model (Zhou et al., 2024b; Wang et al., 2025; Wei et al., 2025). When task identities are unavailable—the task-free setting (Aljundi et al., 2019)—the learner must itself decide *when* to allocate a new module. Published decision rules monitor training-loss spikes or plateaus (Wei et al., 2025; Tathe, 2026), distribution-shift descriptors of intermediate representations (Wang et al., 2025), prospective forgetting after a virtual update (Julian et al., 2025), or the hindrance a new task places on forward learning (Feng et al., 2024). These signals are usually validated by *boundary precision*—whether expansions coincide with the hidden task boundaries of a benchmark—or by ablations against a model that adapts only on the first task.

Boundary precision and the *value* of expansion are different quantities. Most boundaries in standard benchmarks are class-incremental: new labels arrive with new inputs. A learner whose classifier can grow new output rows, on top of a pre-trained representation that already separates the new classes reasonably well, may need no new adapter at all. Whether a new adapter is nevertheless useful can depend on something other than capacity: in exemplar-free learners, class prototypes recorded under an adapter that keeps training become stale (Yu et al., 2020), and freezing that adapter—which is what expansion does—keeps them valid.

This paper asks what common expansion triggers detect, and when what they detect warrants expansion. Our contributions are:

1. **A diagnostic.** We show (Proposition 1) that at a package-level label-novelty event the excess loss of an existing module depends on the input only through its old-label logits, so loss-spike triggers and one-step fresh-module utility respond to classifier novelty, not to representational inadequacy. Empirically, every loss-spike and fresh-utility expansion in our class-incremental streams falls at such an event.
2. **A value-of-expansion protocol.** We measure the value of expansion with oracle expansion at the true boundaries and count-matched random timing, against two references—a single adapter that keeps learning with the same memory budget, and first-session adaptation (Panos et al., 2023; Zhou et al., 2024a)—with trigger thresholds calibrated on shift-free nulls rather than tuned on outcomes.
3. **A regime map.** In the exemplar-based regime, expansion hurts in class-incremental streams yet beats first-session adaptation, so the conclusion of a self-expansion ablation depends on its reference point. In the exemplar-free regime the sign flips [EXT2]. Two mechanism tests support stale classifier statistics as the cause: the flip appears within identical training trajectories when only the prototype rule changes, and exemplar-free drift compensation [EXT4]. The interaction replicates in vision (ViT-B/16, CIFAR-100, ImageNet-R) and inside SEMA’s official implementation [SEMA].
4. **A label-conditional detector, and why detection is not enough.** LABELSURPRISE compares each example’s loss with the recent loss of its own label and ignores labels without history. It never expands at class arrivals and fires when established labels change, but its expansions help only where expansion helps.

The study was pre-registered with frozen code hashes in four stages, uses the seed as the unit of analysis, and reports every failed hypothesis. We do not propose a method that competes with the state of the art; our learner is deliberately simple so that the expansion decision is the only moving part.

2 Related work

Expansion in continual learning. Architectural approaches add capacity per task (Rusu et al., 2016; Yoon et al., 2018) or on demand in task-free streams (Lee et al., 2020; Ye & Bors, 2023). With pre-trained models, adapter pools are expanded per task (Zhou et al., 2024b) or when a shift is detected: SEMA uses autoencoder-based representation descriptors (Wang et al., 2025), Online-LoRA monitors loss plateaus (Wei et al., 2025), LACE expands on loss spikes relative to a recent-loss baseline (Tathe, 2026), CABLE scores prospective forgetting after a virtual update (Julian et al., 2025), and LW2G learns whether to grow from a hindrance metric (Feng et al., 2024). These signals are validated by boundary precision—LACE reports that all of its expansions fall at domain boundaries—or by ablations that remove expansion while keeping first-session adaptation, as in SEMA. We stylise the signals of these families inside one learner, evaluate a LACE-style rule with its published threshold, and run SEMA’s own code with the expansion decision varied.

Reference points for pre-trained-model continual learning. Adapting a pre-trained model on the first session and freezing it is a strong replay-free baseline (Panos et al., 2023; Zhou et al., 2024a), and prototype classifiers on pre-trained features are strong in general (Rebuffi et al., 2017; Mai et al., 2021; McDonnell et al., 2023). Continual fine-tuning of a single model is competitive when representation learning is slowed and the classifier is re-aligned (Zhang et al., 2023). In exemplar-free learning, the features of old classes drift while the model trains; semantic drift compensation estimates this drift from current data (Yu et al., 2020; Gomez-Villa et al., 2024), and SSIAT keeps a single shared adapter with semantic-shift compensation (Tan et al., 2024). We use these as the two references of the value-of-expansion protocol and as an exemplar-free alternative to expansion.

Classifier effects in class-incremental learning. Output-layer bias toward recent classes is well documented (Wu et al., 2019; Hou et al., 2019; Caccia et al., 2022), and the factorisation of class-incremental prediction into task-identity and within-task prediction is established (Kim et al., 2022). We use the corresponding loss identity as a diagnostic of expansion triggers; the identity itself is not new.

Drift and novelty detection in streams. Error-rate drift detectors (Page, 1954; Gama et al., 2004; Bifet & Gavaldà, 2007) and novel-class detection under concept drift (Masud et al., 2011) distinguish concept change from class emergence in classical stream mining, and class-specific monitoring avoids the masking of drift by class-prior changes (Wang et al., 2013). LABELSURPRISE transfers this label-conditional view to the expansion decision.

3 Setting and learner

Stream. A task-free learner observes batches $(\mathbf{x}_t, \mathbf{y}_t)$, $t = 1, \dots, T$, of 16 labelled examples. Segment identities are evaluation metadata and are never shown to the learner, except by the oracle reference policy.

Packages. A frozen encoder f (DistilBERT (Sanh et al., 2019); BERT-base (Devlin et al., 2019) as held-out backbone; ViT-B/16 pre-trained on ImageNet-21k (Dosovitskiy et al., 2021) for vision) is shared. A package k consists of LoRA factors (Hu et al., 2022) on query/value projections (rank 8) and a linear head over the global label space. Only the newest package trains; earlier packages are frozen, as in expand-and-freeze adapter pools (Zhou et al., 2024b; Wang et al., 2025). A new package starts from zero-output LoRA and a zero head.

Training. For every incoming batch the newest package takes two AdamW steps on the batch plus, in the exemplar-based regime, 16 items replayed from its own memory (Chaudhry et al., 2019), with cross-entropy over the labels observed so far.

Memory. A single global reservoir (Vitter, 1985) of 512 examples is shared by the whole pool, so expansion never increases the memory budget.

Two regimes of classifier statistics. Predictions use cosine prototypes; class scores are averaged over packages (equivalently, nearest-class-mean on the normalised concatenation of package features, as in EASE (Zhou et al., 2024b)). In the *exemplar-based* regime, prototypes are recomputed from memory under each package at prediction time, so they always match the current features. In the *exemplar-free* regime (no replay, as in APER, EASE and SEMA), each package accumulates class means of features computed when the examples were observed and never recomputes them; a package spawned after a class was observed holds no prototype for it. Prototypes of a package that keeps training therefore go stale, while those of a frozen package stay valid. Every exemplar-free run also reports the exemplar-based rule on the identical trajectory, which isolates the refresh factor.

Learner competence. A single package reaches 85% on the BANKING77 recurrence stream versus 65% for training-free frozen-feature prototypes (Table 1). An earlier version of this project used a learner that reached only 16%—below the training-free baseline—and we discarded its accuracy comparisons (Appendix A).

4 What a loss signal measures at a label-novelty event

Let S be the labels that may be predicted, partitioned into labels O already trained *by the current package* and labels N in the current batch whose head rows in that package have never been updated. N contains labels new to the stream, but also labels seen before the package was created: in an expand-and-freeze pool, a returning label group is new to the package spawned since it was last seen. We call both *package-level*

label novelty. For an example with label y in group $G(y) \in \{O, N\}$,

$$\ell(x, y) = \underbrace{-\log \sum_{c \in G(y)} p(c | x)}_{\text{label-group term}} - \underbrace{\log p(y | x, G(y))}_{\text{within-group term}}. \quad (1)$$

Proposition 1. *Suppose the head rows of labels in N have never been updated and are zero, so $z_c(x) = 0$ for $c \in N$. For $y \in N$ the loss of an existing package is*

$$\ell(x, y) = \log |N| + \log \left(1 + \frac{1}{|N|} \sum_{o \in O} e^{z_o(x)} \right),$$

where the first term is the within-group term and the second is the label-group term. A fresh package with a zero head incurs $\log(|N| + |O|)$. Hence the advantage of the fresh package before any update (reuse loss minus fresh loss) is $\log((|N| + \sum_o e^{z_o(x)}) / (|N| + |O|))$, which is positive whenever the mean of $e^{z_o(x)}$ over old labels exceeds one, and which depends on x only through the old-label logits.

Proof. With $z_c = 0$ for $c \in N$, $p(y | x) = 1/(|N| + \sum_o e^{z_o})$ and $p(y | x, N) = 1/|N|$; the identity follows from equation 1. A zero head assigns $1/|S|$ to every label. $\square$

The proposition is elementary, but its consequence for expansion triggers is direct. A loss-spike detector at a class arrival measures how confidently the old head predicts old labels; a fresh-versus-reuse utility comparison before or shortly after one update rewards discarding that confident head. Neither asks whether the representation can separate the new labels—a question that a growing head and a prototype classifier answer without new adapter capacity. The label-group term is 36–52% of the boundary loss in our class-incremental streams even after several steps of learning (Figure 1).

4.1 A label-conditional alternative

If label novelty is to be absorbed by the classifier, expansion should respond to evidence that an *established* label’s mapping has changed. Let B_y be the mean of the last 20 batch-mean losses of label y under the current package.

Definition 1 (LABELSURPRISE). *For a batch, $LS = \frac{1}{|\mathcal{E}|} \sum_{i \in \mathcal{E}} (\ell(x_i, y_i) - B_{y_i})$, where $\mathcal{E}$ contains the examples whose label has at least three prior batch-loss records under the current package. LS is undefined (cannot trigger) if $|\mathcal{E}| < 4$.*

Newly introduced labels are excluded by construction, and while a label is being learned its loss falls, so LS is typically negative under class arrivals. Under concept drift, an established label’s loss rises above its own history. LS is an error-based drift statistic (Gama et al., 2004) computed per label. Two earlier variants failed during development (Appendix A).

5 Triggers, references and protocol

Trigger statistics. All triggers are computed on the newest package before it trains on the incoming batch. LABELNOVEL: fraction of batch labels never observed (a trivial reference). LOSSZ: z -score of the incoming loss against a 20-batch window (loss-spike triggers; Tathe, 2026; Wei et al., 2025). REPRZ: z -score of a label-free diagonal-Mahalanobis feature distance (distribution-shift descriptors; Wang et al., 2025). INTERFLOGIT: increase in protected-memory cross-entropy after a virtual update (prospective forgetting; Julian et al., 2025). FRESHUTIL: cross-fitted query-loss advantage of a fresh package over updating the current one after one virtual update (counterfactual utility; Appendix C). Two further statistics, INTERFPROTO and CONFLICTZ, are failed development proposals reported for completeness, and LABELSURPRISE is defined above. Virtual updates use exact snapshot/restore of parameters, optimiser moments and random states; a test verifies bitwise restoration.




Figure 1: Per-batch statistics of a single package on a class-incremental stream (left) and a concept-drift stream (right); development seed 2026. Top: incoming loss split as in equation 1. Middle: loss-spike z -score. Bottom: LABELSURPRISE. Dashed: null-calibrated thresholds ($q = 0.99$). The loss spike crosses its threshold at every class arrival; LABELSURPRISE stays below its threshold under class arrivals and crosses it when established labels flip.

Decision rule and calibration. After the newest package has at least eight updates, a trigger spawns when its statistic exceeds its threshold on two consecutive batches; values in alarm are not added to baseline windows. Each threshold is the 0.99-quantile of the statistic on an i.i.d.-shuffled version of the same stream (same examples, no shift) under the same backbone and regime, using only the development seed. This equalises false-alarm rates across triggers and uses no outcome.

Streams. Class-incremental (CIL): BANKING77 (Casanueva et al., 2020) class-group recurrence (A B A C B) and seven disjoint groups, CLINC150 (Larson et al., 2019) (ten groups of 15), *held-out* 20 Newsgroups (Lang, 1995), Split CIFAR-100 (Krizhevsky, 2009) (ten groups of ten) and ImageNet-R (Hendrycks et al., 2021) (ten groups of 20). Same-label domain shift (DIL): Amazon reviews (Keung et al., 2020) across domains, *held-out* SST-2→Yelp→tweets→IMDB (Socher et al., 2013; Zhang et al., 2015; Barbieri et al., 2020; Maas et al., 2011), and a BREEDS-style (Santurkar et al., 2021) CIFAR-100 superclass stream in which domain d contains the d -th subclass of every superclass. Concept drift: the Amazon and sentiment streams with some domains’ labels flipped, the CIFAR superclass stream with two domains relabelled, and a natural label-semantics stream of TweetEval tasks over the same tweet population. Development examples come from training splits only.

Table 1: Final mean development accuracy (% , prototype-mean inference) over ten seeds in the exemplar-based regime; mean number of packages in parentheses. [†]held-out stream. First-seg.: one package adapted on the first segment, then frozen.

Stream	Frozen	Single	First-seg.	Oracle	Random	LabelNovel	LossZ	ReprZ	InterfLogit	FreshUtil	LS
<i>CIL</i>											
banking_rec	65.0	85.3	69.9	77.1 (5.0)	79.2 (5.0)	81.1 (2.9)	79.5 (4.1)	83.2 (2.8)	83.1 (2.4)	77.3 (5.0)	85.3 (1.0)
banking_cil	46.4	72.7	55.8	61.3 (7.0)	62.8 (7.0)	64.2 (5.1)	62.2 (6.5)	70.5 (1.7)	67.0 (3.3)	62.2 (6.5)	72.7 (1.0)
clinc_cil	58.5	80.8	64.0	70.9 (10.0)	71.4 (10.0)	74.5 (5.2)	71.0 (10.0)	77.8 (2.3)	74.8 (4.9)	71.2 (9.2)	80.8 (1.0)
news_cil [†]	44.7	56.7	48.7	49.5 (5.0)	50.4 (5.0)	56.4 (1.1)	49.4 (5.0)	55.1 (1.5)	49.8 (4.7)	49.4 (5.0)	56.7 (1.0)
<i>DIL</i>											
amazon_rec	81.9	87.0	84.6	86.4 (5.0)	86.9 (5.0)	87.0 (1.0)	87.0 (1.0)	86.6 (1.5)	86.9 (1.2)	87.0 (1.0)	87.0 (1.0)
amazon_dil	82.0	86.3	85.6	86.6 (5.0)	86.5 (5.0)	86.3 (1.0)	86.3 (1.0)	86.1 (2.5)	86.3 (1.0)	86.3 (1.0)	86.3 (1.0)
senti_dil [†]	77.1	81.8	80.1	81.8 (4.0)	81.9 (4.0)	81.8 (1.0)	82.1 (1.6)	82.1 (2.8)	81.9 (1.2)	81.8 (1.0)	81.9 (1.5)
<i>Drift</i>											
amazon_conflict	59.0	60.7	60.7	61.5 (5.0)	61.5 (5.0)	60.7 (1.0)	61.0 (3.2)	61.3 (2.2)	61.4 (1.1)	61.8 (3.8)	61.7 (4.0)
amazon_dilconf	57.0	56.3	57.6	58.2 (5.0)	58.2 (5.0)	56.3 (1.0)	58.3 (5.0)	58.0 (3.1)	57.8 (3.0)	58.4 (4.7)	58.3 (5.0)
senti_conf [†]	55.1	76.3	54.2	58.8 (4.0)	62.2 (4.0)	76.3 (1.0)	61.9 (3.4)	64.4 (2.6)	70.4 (1.9)	63.3 (2.8)	59.6 (4.0)

Policies. *Frozen*: training-free prototypes. *Single*: one package throughout. *First-segment*: one package trained on the first segment only, then frozen, while memory and prototypes keep updating (first-session adaptation). *Oracle*: spawn at every true segment change. *Random*: the oracle’s number of spawns at uniformly random steps. *Triggers*: the statistics above.

Exemplar-free refresh. Semantic drift compensation (SDC; Yu et al., 2020) moves each stored prototype μ_c of the training package by the kernel-weighted mean drift of the current batch, $\Delta_c = \sum_i w_{ic}(z_i^{\text{new}} - z_i^{\text{old}}) / \sum_i w_{ic}$ with $w_{ic} = \exp(-\|z_i^{\text{old}} - \mu_c\|^2 / 2\sigma^2)$, on normalised embeddings with the published $\sigma = 0.3$. We apply it after every update, because a task-free stream has no tasks. SDC changes no training computation and is reported as a further prototype rule on the same trajectories.

Statistics and pre-registration. The study was pre-registered in four stages before any confirmation run of each stage (repository files `notes/tmlr_*prereg.md`, with code hashes): the primary text study (seeds 2027–2036; BERT 2027–2031); a vision, baseline-dependence and natural-drift extension (vision seeds 2027–2034); the exemplar-free regime (text 2027–2036, vision 2027–2034); and ImageNet-R plus drift compensation (5–6 seeds). Development used seed 2026 only. The seed is the unit of analysis. We report paired seed-level mean differences with 95% *t*-intervals and two-sided Wilcoxon signed-rank tests, Holm-corrected within each hypothesis family; deviations are listed in Appendix B.

6 Results

6.1 What the triggers detect

Spawn attribution. In the class-incremental streams without recurrence, every spawn of LOSSZ and FRESHUTIL fell within two batches after a batch containing a never-seen label (BANKING-CIL 55/55 and 55/55, CLINC 81/81 and 74/74, 20News [n/n], CIFAR-100 [n/n]; pooled over seeds). On the recurrence stream BANKING-REC, only 45% and 50% did, so this pre-registered hypothesis failed there. When a label group returns, its labels are new to the package spawned since they were last seen, so Proposition 1 applies; in an analysis declared after the failure, 100% of LOSSZ and FRESHUTIL spawns in every class-incremental stream, BANKING-REC included, fell within two batches of package-level label novelty, against a base rate of 15–30% of eligible batches.

A published threshold. The LACE-style rule (loss above $2.5\times$ its recent mean; Tathe, 2026) [XL].

6.2 The value of expansion with exemplars

Class-incremental: expansion hurts against a learner that keeps training... Oracle expansion at the true boundaries reduced final accuracy relative to a single package by 8.2 points on BANKING-REC, 11.4

on BANKING-CIL, 9.9 on CLINC and 7.1 on held-out 20 Newsgroups (every seed), and [CIFAR] on CIFAR-100 with ViT-B/16. Count-matched random-timing expansion was also harmful, and on BANKING-REC less so than the oracle (79.2% vs. 77.1%).

...but helps against first-session adaptation. The same oracle pools beat the first-session policy by 7.2, 5.5 and 6.9 points (every seed) and by 0.9 points on 20 Newsgroups (n.s.), [CIFAR]. A self-expansion ablation of the usual form would therefore conclude that expansion helps, in exactly the setting where it costs 7–11 points against a single package that keeps learning, which itself beats first-session adaptation by 15–17 points (Figure 2b).

Without rehearsal. Removing replay from the training loss (memory still feeds the prototypes) reduced the harm of expansion: oracle–single was -4.3 on BANKING-REC (significant) and -3.3 , -0.6 , -0.6 on the other text CIL streams (n.s.), so our pre-registered prediction that the harm does not depend on rehearsal was mostly not supported. The difference-in-differences (without minus with replay) was $+3.9$ to $+9.3$ points in every stream (significant).

Domain shift and drift. Under same-label domain shift, oracle expansion was within half a point of the single package. Under concept drift its value was mixed: $+1.9$ points $[0.5, 3.3]$ on Amazon-DILCONF, $+0.8$ $[-2.1, 3.8]$ (n.s.) on Amazon-CONFLICT, and -17.5 $[-20.5, -14.5]$ on the held-out sentiment drift stream, the opposite of our pre-registered prediction. In that stream the domains are distinguishable, a single package trained with replay learns the domain-conditional mapping (82% on the flipped Yelp segment at the end), and frozen packages trained on the unflipped mapping out-vote the correct package under prototype averaging (16%). On the natural label-semantics stream, oracle expansion was -0.8 points [final]. [Vision DIL/drift; proto_own.]

6.3 The value of expansion without exemplars

[EXT2: R1 oracle vs single under proto_ef per CIL stream; R2 within-run DiD; R3 loss_z vs LABELSURPRISE; R4 oracle vs random; firstseg.]

Refreshing without exemplars. [EXT4: S1, S2; sigma sensitivity.]

6.4 A canonical benchmark and a published system

ImageNet-R. [EXT3: I1–I4.]

SEMA. We ran SEMA’s official implementation with its published CIFAR-100 configuration and varied only the expansion decision: the official rule; never expanding (adapters frozen after the first task, the reference of SEMA’s own ablation); never expanding while the adapters keep training on every task; and expanding at every task. [SEMA results: linear head / exemplar-free NCM vs exemplar NCM.]

6.5 Closed-loop triggers and LabelSurprise

Triggers that detect class boundaries spawned 3–5 packages per class-incremental stream and lost 4–10 points relative to the single package in the exemplar-based regime. LABELSURPRISE spawned essentially no package in any class-incremental or domain-shift stream, and so matched the single package (non-inferior in all seven), while on the drift streams it spawned 4–5 packages and tracked the oracle: $+2.1$ $[0.9, 3.3]$ on Amazon-DILCONF, $+1.0$ (n.s.) on Amazon-CONFLICT and -16.7 on the held-out sentiment drift stream, where expansion itself is harmful. In the exemplar-free regime [R3]. [BERT, vision triggers.]




Figure 2: Regime map. Paired seed-level difference in final accuracy (mean, 95% t -interval) between oracle expansion and (a) a single package that keeps learning, (b) first-segment adaptation, with exemplar-refreshed prototypes (blue) and exemplar-free prototypes (orange). [†]held-out stream.

7 Discussion and limitations

What expansion buys. [Synthesis: expansion at classifier novelty freezes the features under which classifier statistics were recorded; it pays when those statistics cannot be refreshed and costs when they can. First-session adaptation, drift compensation and exemplar refresh are competing remedies.]

Scope. Our encoders are small to medium, our streams are short (80–500 batches), and three of five drift streams use synthetic relabelling. Our triggers are stylised versions of published signals inside a common learner; the SEMA experiment checks one complete system, and we make no claim about reported benchmark numbers. The value of expansion also depends on the inference rule: routing each input to the package whose own prototypes fit best can recover part of the benefit of expansion under domain shift in vision (Appendix tables report all rules).

Calibration. Thresholds are set on an i.i.d.-shuffled version of the same stream. This equalises false-alarm rates across triggers, which a comparison of triggers requires, but it is not available to a deployed learner. The conclusion that loss-type triggers fire at label arrivals does not depend on the threshold, since their AUROC for class arrivals is near one.

What LabelSurprise detects. LABELSURPRISE responds to an increase in the loss of established labels. Concept drift produces this, but so does a within-label shift of $P(x | y)$, such as a new subpopulation under an existing superclass. It is therefore a detector of *established-label degradation*, not of concept drift per se.

Training budget of the references. Our first-segment policy trains for the same number of steps per batch as every other policy; SEMA and APER adapt for full epochs on the first task. The SEMA experiment uses SEMA’s own budget.

Recommendations. When an expansion trigger is proposed, its evaluation should include (i) the measured value of expansion in the evaluated regime against a single module that keeps learning with the same memory budget, and against first-session adaptation; (ii) a statement of whether classifier statistics are refreshed, since this decides the sign of the value of expansion at class arrivals; (iii) separation of label novelty from changes in established labels, since boundary precision on class-incremental benchmarks is achieved by label-novelty detection alone; and (iv) thresholds calibrated on shift-free nulls rather than tuned on outcomes.

8 Conclusion

[TO BE WRITTEN AFTER CONFIRMATION.]

Use of AI assistants. [Statement: which parts were assisted by an AI coding/writing assistant; all claims, code and text were checked by the authors, who take full responsibility.]

References

- Rahaf Aljundi, Klaas Kelchtermans, and Tinne Tuytelaars. Task-free continual learning. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2019.
- Francesco Barbieri, Jose Camacho-Collados, Luis Espinosa Anke, and Leonardo Neves. TweetEval: Unified benchmark and comparative evaluation for tweet classification. In *Findings of EMNLP*, 2020.
- Albert Bifet and Ricard Gavaldà. Learning from time-changing data with adaptive windowing. In *Proceedings of the SIAM International Conference on Data Mining*, 2007.
- Lucas Caccia, Rahaf Aljundi, Nader Asadi, Tinne Tuytelaars, Joelle Pineau, and Eugene Belilovsky. New insights on reducing abrupt representation change in online continual learning. In *International Conference on Learning Representations*, 2022.
- Iñigo Casanueva, Tadas Temčinas, Daniela Gerz, Matthew Henderson, and Ivan Vulić. Efficient intent detection with dual sentence encoders. In *Proceedings of the 2nd Workshop on Natural Language Processing for Conversational AI*, 2020.
- Arslan Chaudhry, Marcus Rohrbach, Mohamed Elhoseiny, Thalaiyasingam Ajanthan, Puneet K. Dokania, Philip H. S. Torr, and Marc’Aurelio Ranzato. On tiny episodic memories in continual learning. *arXiv preprint arXiv:1902.10486*, 2019.
- Jacob Devlin, Ming-Wei Chang, Kenton Lee, and Kristina Toutanova. BERT: Pre-training of deep bidirectional transformers for language understanding. In *Proceedings of NAACL-HLT*, 2019.
- Alexey Dosovitskiy, Lucas Beyer, Alexander Kolesnikov, Dirk Weissenborn, Xiaohua Zhai, Thomas Unterthiner, Mostafa Dehghani, Matthias Minderer, Georg Heigold, Sylvain Gelly, Jakob Uszkoreit, and Neil Houlsby. An image is worth 16x16 words: Transformers for image recognition at scale. In *International Conference on Learning Representations*, 2021.
- Qian Feng, Da-Wei Zhou, Hanbin Zhao, Chao Zhang, Jiahua Dong, Dengxin Dai, and Hui Qian. LW2G: Learning whether to grow for prompt-based continual learning. *arXiv preprint arXiv:2409.18860*, 2024.
- João Gama, Pedro Medas, Gladys Castillo, and Pedro Rodrigues. Learning with drift detection. In *Brazilian Symposium on Artificial Intelligence (SBIA)*, 2004.
- Alex Gomez-Villa, Dipam Goswami, Kai Wang, Andrew D. Bagdanov, Bartłomiej Twardowski, and Joost van de Weijer. Exemplar-free continual representation learning via learnable drift compensation. In *Proceedings of the European Conference on Computer Vision (ECCV)*, 2024.
- Dan Hendrycks, Steven Basart, Norman Mu, Saurav Kadavath, Frank Wang, Evan Dorundo, Rahul Desai, Tyler Zhu, Samyak Parajuli, Mike Guo, Dawn Song, Jacob Steinhardt, and Justin Gilmer. The many faces of robustness: A critical analysis of out-of-distribution generalization. In *Proceedings of the IEEE/CVF International Conference on Computer Vision (ICCV)*, 2021.

-
- Saihui Hou, Xinyu Pan, Chen Change Loy, Zilei Wang, and Dahua Lin. Learning a unified classifier incrementally via rebalancing. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2019.
- Edward J. Hu, Yelong Shen, Phillip Wallis, Zeyuan Allen-Zhu, Yuanzhi Li, Shean Wang, Lu Wang, and Weizhu Chen. LoRA: Low-rank adaptation of large language models. In *International Conference on Learning Representations*, 2022.
- Jack Julian, Yun Sing Koh, and Albert Bifet. Building adaptive knowledge bases for evolving continual learning models. *npj Artificial Intelligence*, 1:26, 2025. doi: 10.1038/s44387-025-00028-4.
- Phillip Keung, Yichao Lu, György Szarvas, and Noah A. Smith. The multilingual Amazon reviews corpus. In *Proceedings of EMNLP*, 2020.
- Gyuhak Kim, Changnan Xiao, Tatsuya Konishi, Zixuan Ke, and Bing Liu. A theoretical study on solving continual learning. In *Advances in Neural Information Processing Systems*, 2022.
- Alex Krizhevsky. Learning multiple layers of features from tiny images. Technical report, University of Toronto, 2009.
- Ken Lang. Newsweeder: Learning to filter netnews. In *Proceedings of ICML*, 1995.
- Stefan Larson, Anish Mahendran, Joseph J. Peper, Christopher Clarke, Andrew Lee, Parker Hill, Jonathan K. Kummerfeld, Kevin Leach, Michael A. Laurenzano, Lingjia Tang, and Jason Mars. An evaluation dataset for intent classification and out-of-scope prediction. In *Proceedings of EMNLP-IJCNLP*, 2019.
- Soochan Lee, Junsoo Ha, Dongsu Zhang, and Gunhee Kim. A neural Dirichlet process mixture model for task-free continual learning. In *International Conference on Learning Representations*, 2020.
- Andrew L. Maas, Raymond E. Daly, Peter T. Pham, Dan Huang, Andrew Y. Ng, and Christopher Potts. Learning word vectors for sentiment analysis. In *Proceedings of ACL*, 2011.
- Zheda Mai, Ruiwen Li, Hyunwoo Kim, and Scott Sanner. Supervised contrastive replay: Revisiting the nearest class mean classifier in online class-incremental continual learning. In *CVPR Workshops*, 2021.
- Mohammad M. Masud, Jing Gao, Latifur Khan, Jiawei Han, and Bhavani Thuraisingham. Classification and novel class detection in concept-drifting data streams under time constraints. *IEEE Transactions on Knowledge and Data Engineering*, 23(6):859–874, 2011.
- Mark D. McDonnell, Dong Gong, Amin Parvaneh, Ehsan Abbasnejad, and Anton van den Hengel. RanPAC: Random projections and pre-trained models for continual learning. In *Advances in Neural Information Processing Systems*, 2023.
- E. S. Page. Continuous inspection schemes. *Biometrika*, 41(1/2):100–115, 1954.
- Aristeidis Panos, Yuriko Kobe, Daniel Olmeda Reino, Rahaf Aljundi, and Richard E. Turner. First session adaptation: A strong replay-free baseline for class-incremental learning. In *Proceedings of the IEEE/CVF International Conference on Computer Vision (ICCV)*, pp. 18820–18830, 2023.
- Sylvestre-Alvise Rebuffi, Alexander Kolesnikov, Georg Sperl, and Christoph H. Lampert. iCaRL: Incremental classifier and representation learning. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, 2017.
- Andrei A. Rusu, Neil C. Rabinowitz, Guillaume Desjardins, Hubert Soyer, James Kirkpatrick, Koray Kavukcuoglu, Razvan Pascanu, and Raia Hadsell. Progressive neural networks. *arXiv preprint arXiv:1606.04671*, 2016.
- Victor Sanh, Lysandre Debut, Julien Chaumond, and Thomas Wolf. DistilBERT, a distilled version of BERT: smaller, faster, cheaper and lighter. *arXiv preprint arXiv:1910.01108*, 2019.

-
- Shibani Santurkar, Dimitris Tsipras, and Aleksander Madry. BREEDS: Benchmarks for subpopulation shift. In *International Conference on Learning Representations*, 2021.
- Richard Socher, Alex Perelygin, Jean Wu, Jason Chuang, Christopher D. Manning, Andrew Ng, and Christopher Potts. Recursive deep models for semantic compositionality over a sentiment treebank. In *Proceedings of EMNLP*, 2013.
- Yuwen Tan, Qinzhao Zhou, Xiang Xiang, Ke Wang, Yuchuan Wu, and Yongbin Li. Semantically-shifted incremental adapter-tuning is a continual ViTransformer. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pp. 23252–23262, 2024.
- Shivnath Tathe. LACE: Loss-adaptive capacity expansion for continual learning. *arXiv preprint arXiv:2603.28611*, 2026.
- Jeffrey S. Vitter. Random sampling with a reservoir. *ACM Transactions on Mathematical Software*, 11(1): 37–57, 1985.
- Huiyi Wang, Haodong Lu, Lina Yao, and Dong Gong. Self-expansion of pre-trained models with mixture of adapters for continual learning. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2025.
- Shuo Wang, Leandro L. Minku, Davide Ghezzi, Daniele Caltabiano, Peter Tino, and Xin Yao. Concept drift detection for online class imbalance learning. In *International Joint Conference on Neural Networks (IJCNN)*, 2013.
- Xiwen Wei, Guihong Li, and Radu Marculescu. Online-LoRA: Task-free online continual learning via low rank adaptation. In *Proceedings of the IEEE/CVF Winter Conference on Applications of Computer Vision (WACV)*, 2025.
- Yue Wu, Yinpeng Chen, Lijuan Wang, Yuancheng Ye, Zicheng Liu, Yandong Guo, and Yun Fu. Large scale incremental learning. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2019.
- Fei Ye and Adrian G. Bors. Self-evolved dynamic expansion model for task-free continual learning. In *Proceedings of the IEEE/CVF International Conference on Computer Vision (ICCV)*, 2023.
- Jaehong Yoon, Eunho Yang, Jeongtae Lee, and Sung Ju Hwang. Lifelong learning with dynamically expandable networks. In *International Conference on Learning Representations*, 2018.
- Lu Yu, Bartłomiej Twardowski, Xialei Liu, Luis Herranz, Kai Wang, Yongmei Cheng, Shangling Jui, and Joost van de Weijer. Semantic drift compensation for class-incremental learning. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pp. 6982–6991, 2020.
- Gengwei Zhang, Liyuan Wang, Guoliang Kang, Ling Chen, and Yunchao Wei. SLCA: Slow learner with classifier alignment for continual learning on a pre-trained model. In *Proceedings of the IEEE/CVF International Conference on Computer Vision (ICCV)*, pp. 19148–19158, 2023.
- Xiang Zhang, Junbo Zhao, and Yann LeCun. Character-level convolutional networks for text classification. In *Advances in Neural Information Processing Systems*, 2015.
- Da-Wei Zhou, Zi-Wen Cai, Han-Jia Ye, De-Chuan Zhan, and Ziwei Liu. Revisiting class-incremental learning with pre-trained models: Generalizability and adaptivity are all you need. *International Journal of Computer Vision*, 2024a.
- Da-Wei Zhou, Hai-Long Sun, Han-Jia Ye, and De-Chuan Zhan. Expandable subspace ensemble for pre-trained model-based class-incremental learning. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2024b.

A Development history and failed candidates

This project began as the design of a task-free allocator (Appendix C). Its learner took a single AdamW step per 16-example batch into a 77-way head and reached about 16% final accuracy on the BANKING recurrence stream with label-free routing (22% with a learned router), whereas training-free frozen-feature prototypes reached 71% (logistic probe on frozen features: 86%). We therefore discarded all of its accuracy comparisons and built the learner of Section 3. All development of the new testbed used seed 2026 only:

1. INTERPROTO (head-equalised prospective interference) was our first proposal. On the seed-2026 shadow traces it did not detect concept drift (AUROC 0.54 and 0.22 on the two development drift streams).
2. CONFLICTZ v1 restricted the loss z -score to examples with previously observed labels. It detected class-incremental boundaries (AUROC 1.00) because labels observed one batch earlier are still untrained. v2 required labels to be first observed at least eight batches earlier; it still fired in class-incremental streams, because newly learned labels have higher loss than long-established ones. This motivated the label-conditional baseline of LABELSURPRISE.
3. The first closed-loop pilot added alarmed values to the z -score baseline windows, so a genuine spike inflated its own baseline. Excluding alarmed values applies identically to all z -type triggers.
4. A slower single learner (learning rate 10^{-4}) did not reduce exemplar-free staleness in development (BANKING-CIL: 13.1% vs. 26.2% at 10^{-3}) and was not carried into the confirmation.

The thresholds, learner hyperparameters (learning rate 10^{-3} , two steps per batch, replay 16, memory 512) and all trigger constants were fixed before the confirmation study.

B Pre-registration deviations

[List from notes: $q=0.95$ grid not run; BERT subset; H1 failure and post-hoc package-level analysis; A1 addendum; ext2 duplicate-launch note; reduced seeds in extensions 3–4.]

C Case study: an exact counterfactual allocator

Before this study, we developed and audited an allocator, CAU (counterfactual action utility), that compares an exact, optimiser-faithful one-step update of each existing package and of a fresh classifier-bearing package against a common no-update reference on cross-fitted query halves, subject to a separate protected-memory feasibility check, and spawns after two consecutive windows of positive fresh utility. On a BANKING77 class-group recurrence stream (A B A C B), CAU spawned exactly at the arrivals of the B and C label groups in all ten seed/order development trajectories. A *head-only* variant—identical, except that LoRA outputs are held at zero and only the classifier trains—made identical allocation decisions in all ten trajectories. At the original spawn windows, about 87% of the fresh-versus-best-reuse utility advantage was already present before the prospective update, i.e. it reflected the replacement of the old classifier rather than learning. These observations, made with the weaker learner of Appendix A and reported only for decisions, motivated Proposition 1 and the present study.